

CNN on MNIST

```
In [2]: from future import print_function
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Conv2D, MaxPool2D, Flatten, Dropout, BatchNormalization
from keras import backend as K
import matplotlib.pyplot as plt

Using TensorFlow backend.
```

```
In [3]: %matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6ee59eaf1d90fcb3b534a2562ccea4
# https://stackoverflow.com/answer/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, color=['b']):
    ax.plot(x, vy, 'b', label='Validation Loss')
    ax.plot(x, ty, 'r', label='Train Loss')
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

```
In [4]: batch_size = 128
num_classes = 10
epochs = 12

# input image dimensions
img_rows, img_cols = 28, 28

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

```
In [5]: print(x_train.data_format())
print(x_train.shape)

channels_last
(60000, 28, 28)
```

```
In [6]: x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
input_shape = (1, img_rows, img_cols)
```

```
In [7]: x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

x_train shape: (60000, 1, 28, 28)
60000 train samples
10000 test samples
```

CNN Model of 3 Layers With BatchNormalization, Dropout and MaxPooling

```
In [37]: model3= Sequential()

model3.add(Conv2D(32, kernel_size=3, activation='relu', input_shape=input_shape, data_format='channels_first'))
model3.add(BatchNormalization(axis=1)) #for conv layer axis=1 ,,

model3.add(Conv2D(160, kernel_size=4, activation='relu'))
model3.add(BatchNormalization(axis=1)) #for conv layer axis=1
model3.add(Dropout(0.3))
model3.add(MaxPool2D(pool_size=2))

model3.add(Conv2D(196, kernel_size=3, activation='relu'))
model3.add(BatchNormalization(axis=1)) #for conv layer axis=1

model3.add(Flatten())
model3.add(Dense(128, activation='relu'))
model3.add(Dense(64, activation='relu'))
model3.add(Dropout(0.5))
model3.add(Dense(num_classes, activation='softmax'))

In [38]: model3.summary()
```

Layer (type)	Output Shape	Param #
conv2d_36 (Conv2D)	(None, 32, 26, 26)	320
batch_normalization_27 (Batch Normalization)	(None, 32, 26, 26)	104
conv2d_31 (Conv2D)	(None, 29, 23, 160)	66720
batch_normalization_28 (Batch Normalization)	(None, 29, 23, 160)	640
dropout_23 (Dropout)	(None, 29, 23, 160)	0
max_pooling2d_8 (MaxPooling2D)	(None, 14, 11, 160)	0
conv2d_32 (Conv2D)	(None, 12, 9, 196)	282436
batch_normalization_29 (Batch Normalization)	(None, 12, 9, 196)	784
flatten_5 (Flatten)	(None, 21168)	0
dense_13 (Dense)	(None, 128)	2709632
dense_14 (Dense)	(None, 64)	8256
dropout_24 (Dropout)	(None, 64)	0
dense_15 (Dense)	(None, 10)	650
Total params: 3,069,542		
Trainable params: 3,068,778		
Non-trainable params: 764		

```
In [39]: model3.compile(loss=keras.losses.categorical_crossentropy,
                    optimizer=keras.optimizers.Adadelta(),
                    metrics=['accuracy'])

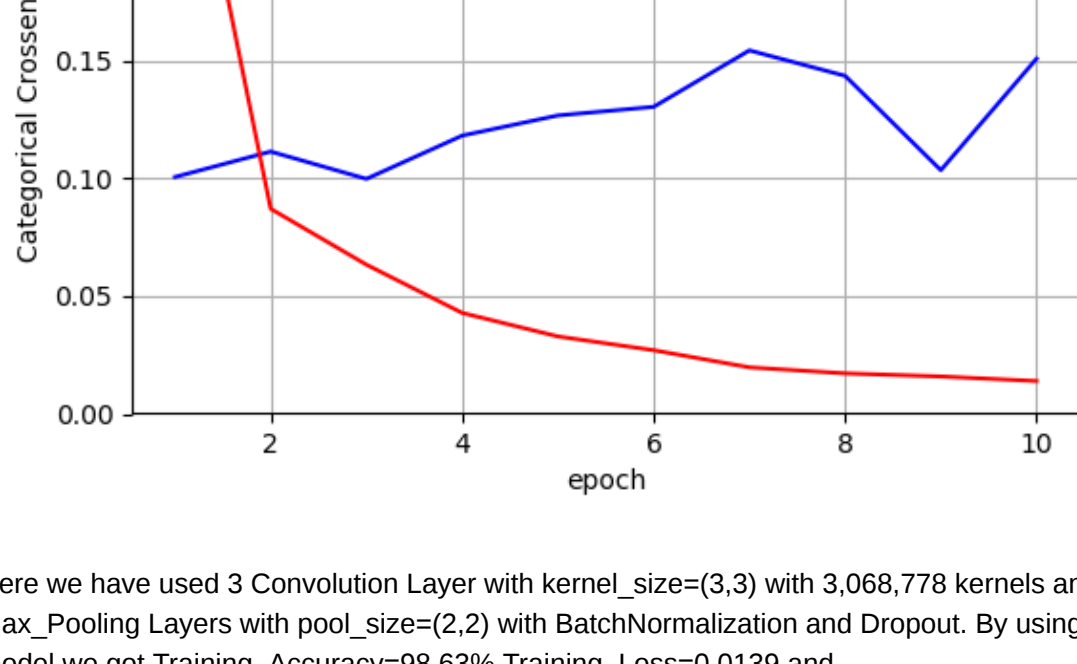
history3= model3.fit(x_train, y_train,
                    batch_size=128,
                    epochs=10,
                    verbose=1,
                    validation_data=(x_test, y_test))
score = model3.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

Train on 60000 samples, validate on 10000 samples
Epoch 1/10
60000/60000 [=====] - 266s 4ms/step - loss: 0.2863 - acc: 0.9235 - val_loss: 0.1805 - val_acc: 0.9764
Epoch 2/10
60000/60000 [=====] - 261s 4ms/step - loss: 0.0878 - acc: 0.9789 - val_loss: 0.1113 - val_acc: 0.9788
Epoch 3/10
60000/60000 [=====] - 261s 4ms/step - loss: 0.0633 - acc: 0.9849 - val_loss: 0.0997 - val_acc: 0.9859
Epoch 4/10
60000/60000 [=====] - 261s 4ms/step - loss: 0.0428 - acc: 0.9893 - val_loss: 0.1181 - val_acc: 0.9848
Epoch 5/10
60000/60000 [=====] - 261s 4ms/step - loss: 0.0329 - acc: 0.9922 - val_loss: 0.1266 - val_acc: 0.9863
Epoch 6/10
60000/60000 [=====] - 261s 4ms/step - loss: 0.0278 - acc: 0.9935 - val_loss: 0.1383 - val_acc: 0.9856
Epoch 7/10
60000/60000 [=====] - 261s 4ms/step - loss: 0.0197 - acc: 0.9954 - val_loss: 0.1542 - val_acc: 0.9854
Epoch 8/10
60000/60000 [=====] - 261s 4ms/step - loss: 0.0172 - acc: 0.9961 - val_loss: 0.1434 - val_acc: 0.9860
Epoch 9/10
60000/60000 [=====] - 261s 4ms/step - loss: 0.0159 - acc: 0.9961 - val_loss: 0.1034 - val_acc: 0.9897
Epoch 10/10
60000/60000 [=====] - 261s 4ms/step - loss: 0.0139 - acc: 0.9969 - val_loss: 0.1508 - val_acc: 0.9863
Test loss: 0.15075194108878912
Test accuracy: 0.9863
```

```
In [40]: score= model3.evaluate(x_test, y_test, verbose=0)
print('Test score: ',score[0])
print('Test accuracy: ',score[1])

Test score: 0.15075194108878912
Test accuracy: 0.9863
```

```
In [41]: fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
x = list(range(1,11))
vy = model3.history.history['val_loss']
ty = model3.history.history['loss']
plt_dynamic(x, vy, ty, ax)
```



Here we have used 3 Convolution Layer with kernel_size=(3,3) with 3,068,778 kernels and 1 Max_Pooling Layers with pool_size=(2,2) with BatchNormalization and Dropout. By using this model we got Training_Accuracy=98.63%, Training_Loss=0.0139 and Test_Accuracy=98.63%, Test_Loss=0.1507

CNN Model of 4 Layers With BatchNormalization, Dropout and MaxPooling Layer.

```
In [9]: model= Sequential()

model.add(Conv2D(32, kernel_size=(3,3), activation='relu', input_shape=input_shape, data_format='channels_first'))
model.add(BatchNormalization(axis=1)) #for conv layer axis=1 ,,

model.add(Conv2D(35, kernel_size=3, activation='relu'))
model.add(BatchNormalization(axis=1)) #for conv layer axis=1
model.add(MaxPool2D(pool_size=2))
model.add(Dropout(0.25))

model.add(Conv2D(70, kernel_size=3, activation='relu'))
model.add(BatchNormalization(axis=1)) #for conv layer axis=1
model.add(Dropout(0.25))

model.add(Conv2D(100, kernel_size=2, activation='relu'))
model.add(BatchNormalization(axis=1)) #for conv layer axis=1
model.add(MaxPool2D(pool_size=2))

model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dense(64, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

WARNING:tensorflow:From C:\Anaconda3\lib\site-packages\tensorflow\python\tensorflow\ops\py_op_def_library.py:263: colocate_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.
Instructions for updating:
Colocations handled automatically by placer.
WARNING:tensorflow:From C:\Anaconda3\lib\site-packages\keras\backend\tensorflow_backend.py:3445: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.
Instructions for updating:
Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

In [10]: model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 32, 26, 26)	320
batch_normalization_1 (Batch Normalization)	(None, 32, 26, 26)	128
conv2d_2 (Conv2D)	(None, 30, 24, 35)	8225
batch_normalization_2 (Batch Normalization)	(None, 30, 24, 35)	120
max_pooling2d_1 (MaxPooling2D)	(None, 15, 12, 35)	0
dropout_1 (Dropout)	(None, 15, 12, 35)	0
conv2d_3 (Conv2D)	(None, 13, 10, 70)	22120
batch_normalization_3 (Batch Normalization)	(None, 13, 10, 70)	52
dropout_2 (Dropout)	(None, 13, 10, 70)	0
conv2d_4 (Conv2D)	(None, 12, 9, 100)	28100
batch_normalization_4 (Batch Normalization)	(None, 12, 9, 100)	48
max_pooling2d_2 (MaxPooling2D)	(None, 6, 4, 100)	0
flatten_1 (Flatten)	(None, 2400)	0
dense_1 (Dense)	(None, 128)	307328
dense_2 (Dense)	(None, 64)	8256
dropout_3 (Dropout)	(None, 64)	0
dense_3 (Dense)	(None, 10)	650
Total params: 375,347		
Trainable params: 375,173		
Non-trainable params: 174		

```
In [11]: model.compile(loss=keras.losses.categorical_crossentropy,
                    optimizer=keras.optimizers.Adadelta(),
                    metrics=['accuracy'])


model.fit(x_train, y_train,
        batch_size=batch_size,
        epochs=epochs,
        verbose=1,
        validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

WARNING:tensorflow:From C:\Anaconda3\lib\site-packages\tensorflow\python\tensorflow\ops\py_op_def_library.py:3866: int32 (from tensorflow.python.ops.math_ops) is deprecated and will be removed in a future version.
Instructions for updating:
Use tf.cast instead.
Train on 60000 samples, validate on 10000 samples
Epoch 1/12
60000/60000 [=====] - 384s 6ms/step - loss: 0.4619 - acc: 0.8613 - val_loss: 0.1801 - val_acc: 0.9718
Epoch 2/12
60000/60000 [=====] - 365s 6ms/step - loss: 0.1266 - acc: 0.9663 - val_loss: 0.1029 - val_acc: 0.9726
Epoch 3/12
60000/60000 [=====] - 364s 6ms/step - loss: 0.0838 - acc: 0.9779 - val_loss: 0.0835 - val_acc: 0.9867
Epoch 4/12
60000/60000 [=====] - 381s 6ms/step - loss: 0.0639 - acc: 0.9829 - val_loss: 0.0476 - val_acc: 0.9885
Epoch 5/12
60000/60000 [=====] - 378s 6ms/step - loss: 0.0562 - acc: 0.9859 - val_loss: 0.0389 - val_acc: 0.9889
Epoch 6/12
60000/60000 [=====] - 381s 6ms/step - loss: 0.0452 - acc: 0.9877 - val_loss: 0.0512 - val_acc: 0.9897
Epoch 7/12
60000/60000 [=====] - 367s 6ms/step - loss: 0.0429 - acc: 0.9885 - val_loss: 0.0378 - val_acc: 0.9905
Epoch 8/12
60000/60000 [=====] - 367s 6ms/step - loss: 0.0367 - acc: 0.9900 - val_loss: 0.0420 - val_acc: 0.9916
Epoch 9/12
60000/60000 [=====] - 368s 6ms/step - loss: 0.0335 - acc: 0.9910 - val_loss: 0.0379 - val_acc: 0.9922
Epoch 10/12
60000/60000 [=====] - 367s 6ms/step - loss: 0.0298 - acc: 0.9922 - val_loss: 0.0374 - val_acc: 0.9918
Epoch 11/12
60000/60000 [=====] - 366s 6ms/step - loss: 0.0256 - acc: 0.9934 - val_loss: 0.0375 - val_acc: 0.9924
Epoch 12/12
60000/60000 [=====] - 378s 6ms/step - loss: 0.0241 - acc: 0.9938 - val_loss: 0.0361 - val_acc: 0.9925
Test loss: 0.03613165665459746
Test accuracy: 0.9925
```

```
In [12]: score= model.evaluate(x_test, y_test, verbose=0)
print('Test score: ',score[0])
print('Test accuracy: ',score[1])

Test score: 0.03613165665459746
Test accuracy: 0.9925
```

```
In [13]: fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
x = list(range(1,13))
vy = model.history.history['val_loss']
ty = model.history.history['loss']
plt_dynamic(x, vy, ty, ax)
```



Here we have used 4 Convolution Layer with kernel_size=(3,3) with 375,173 kernels and 2 Max_Pooling Layers with pool_size=(2,2) with BatchNormalization and Dropout. By using this model we got Training_Accuracy=99.38%, Training_Loss=0.0241 and Test_Accuracy=99.25%, Test_Loss=0.0361

CNN Model of 7 Layers With BatchNormalization, Dropout and MaxPooling Layer.

```
In [20]: model2.reset_states()

In [21]: model2= Sequential()

model2.add(Conv2D(32, kernel_size=6, activation='relu', input_shape=input_shape, data_format='channels_first'))
model2.add(BatchNormalization(axis=1)) #for conv layer axis=1 ,,

model2.add(Conv2D(150, kernel_size=5, activation='relu'))
model2.add(BatchNormalization(axis=1)) #for conv layer axis=1
model2.add(Dropout(0.3))

model2.add(Conv2D(150, kernel_size=4, activation='relu'))
model2.add(BatchNormalization(axis=1)) #for conv layer axis=1
model2.add(Dropout(0.2))

model2.add(Conv2D(200, kernel_size=4, padding='same', activation='relu'))
model2.add(BatchNormalization(axis=1)) #for conv layer axis=1
model2.add(Dropout(0.25))

model2.add(Conv2D(200, kernel_size=3, activation='relu'))
model2.add(BatchNormalization(axis=1)) #for conv layer axis=1
model2.add(Dropout(0.35))

model2.add(Conv2D(150, kernel_size=3, padding='same', activation='relu'))
model2.add(Dropout(0.25))

model2.add(Conv2D(180, kernel_size=3, activation='relu'))
model2.add(BatchNormalization(axis=1)) #for conv layer axis=1
model2.add(MaxPool2D(pool_size=2))

model2.add(Flatten())
model2.add(Dense(128, activation='relu'))
model2.add(Dense(64, activation='relu'))
model2.add(Dropout(0.5))
model2.add(Dense(num_classes, activation='softmax'))

In [22]: model2.summary()
```

Layer (type)	Output Shape	Param #
conv2d_28 (Conv2D)	(None, 32, 23, 23)	1184
batch_normalization_18 (Batch Normalization)	(None, 32, 23, 23)	92
conv2d_21 (Conv2D)	(None, 28, 19, 150)	86400
batch_normalization_19 (Batch Normalization)	(None, 28, 19, 150)	600
dropout_15 (Dropout)	(None, 28, 19, 150)	0
conv2d_22 (Conv2D)	(None, 25, 16, 150)	360150
batch_normalization_20 (Batch Normalization)	(None, 25, 16, 150)	600
dropout_16 (Dropout)	(None, 25, 16, 150)	0
conv2d_23 (Conv2D)	(None, 25, 16, 200)	480200
batch_normalization_21 (Batch Normalization)	(None, 25, 16, 200)	800
dropout_17 (Dropout)	(None, 25, 16, 200)	0
conv2d_24 (Conv2D)	(None, 23, 14, 200)	360200
batch_normalization_22 (Batch Normalization)	(None, 23, 14, 200)	800
dropout_18 (Dropout)	(None, 23, 14, 200)	0
conv2d_25 (Conv2D)	(None, 23, 14, 150)	270150
dropout_19 (Dropout)	(None, 23, 14, 150)	0
conv2d_26 (Conv2D)	(None, 21, 12, 180)	243180
batch_normalization_23 (Batch Normalization)	(None, 21, 12, 180)	720
max_pooling2d_5 (MaxPooling2D)	(None, 10, 6, 180)	0
flatten_3 (Flatten)	(None, 10800)	0
dense_7 (Dense)	(None, 128)	1382528
dense_8 (Dense)	(None, 64)	8256
dropout_20 (Dropout)	(None, 64)	0
dense_9 (Dense)	(None, 10)	650
Total params: 3,196,510		
Trainable params: 3,194,704		
Non-trainable params: 1,806		

```
In [23]: model2.compile(loss=keras.losses.categorical_crossentropy,
                    optimizer=keras.optimizers.Adadelta(),
                    metrics=['accuracy'])


history2= model2.fit(x_train, y_train,
                    batch_size=150,
                    epochs=13,
                    verbose=1,
                    validation_data=(x_test, y_test))
score = model2.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

Train on 60000 samples, validate on 10000 samples
Epoch 1/13
60000/60000 [=====] - 902s 15ms/step - loss: 0.5966 - acc: 0.8275 - val_loss: 0.2085 - val_acc: 0.9585
Epoch 2/13
60000/60000 [=====] - 888s 15ms/step - loss: 0.1367 - acc: 0.9654 - val_loss: 0.0849 - val_acc: 0.9809
Epoch 3/13
60000/60000 [=====] - 888s 15ms/step - loss: 0.0893 - acc: 0.9787 - val_loss: 0.0911 - val_acc: 0.9822
Epoch 4/13
60000/60000 [=====] - 887s 15ms/step - loss: 0.0714 - acc: 0.9831 - val_loss: 0.0464 - val_acc: 0.9907
Epoch 5/13
60000/60000 [=====] - 887s 15ms/step - loss: 0.0595 - acc: 0.9858 - val_loss: 0.0490 - val_acc: 0.9900
Epoch 6/13
60000/60000 [=====] - 887s 15ms/step - loss: 0.0493 - acc: 0.9876 - val_loss: 0.0521 - val_acc: 0.9900
Epoch 7/13
60000/60000 [=====] - 887s 15ms/step - loss: 0.0418 - acc: 0.9897 - val_loss: 0.0526 - val_acc: 0.9908
Epoch 8/13
60000/60000 [=====] - 888s 15ms/step - loss: 0.0368 - acc: 0.9907 - val_loss: 0.0540 - val_acc: 0.9905
Epoch 9/13
60000/60000 [=====] - 887s 15ms/step - loss: 0.0345 - acc: 0.9915 - val_loss: 0.0572 - val_acc: 0.9907
Epoch 10/13
60000/60000 [=====] - 887s 15ms/step - loss: 0.0244 - acc: 0.9921 - val_loss: 0.0551 - val_acc: 0.9907
Epoch 11/13
60000/60000 [=====] - 887s 15ms/step - loss: 0.0300 - acc: 0.9929 - val_loss: 0.0758 - val_acc: 0.9880
Epoch 12/13
60000/60000 [=====] - 887s 15ms/step - loss: 0.0244 - acc: 0.9944 - val_loss: 0.0449 - val_acc: 0.9906
Epoch 13/13
60000/60000 [=====] - 888s 15ms/step - loss: 0.0263 - acc: 0.9940 - val_loss: 0.0555 - val_acc: 0.9909
Test loss: 0.05553952941512905
Test accuracy: 0.9909
```

```
In [24]: score= model2.evaluate(x_test, y_test, verbose=0)
print('Test score: ',score[0])
print('Test accuracy: ',score[1])

Test score: 0.05553952941512905
Test accuracy: 0.9909
```

```
In [25]: fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
x = list(range(1,14))
vy = model2.history.history['val_loss']
ty = model2.history.history['loss']
plt_dynamic(x, vy, ty, ax)
```



Here we have used 7 Convolution Layer with kernel_size=(3,3) with 3,194,704 kernels and 1 Max_Pooling Layers with pool_size=(2,2) with BatchNormalization and Dropout. By using this model we got Training_Accuracy=99.40%, Training_Loss=0.0263 and Test_Accuracy=99.09%, Test_Loss=0.0555

```
In [8]: from prettytable import PrettyTable
x = PrettyTable()
x.field_names = [ "No.Convolution_Layers", "Kernel_Sizes", "No.of_Max_Pooling_Layers", "Pool_Size", "Test_Accuracy" ]

In [9]: x.add_row(['3','(3x3)','1','(2x2)','98.63%'])
x.add_row(['4','(3x3)','2','(2x2)','99.35%'])
x.add_row(['7','(3x3)','1','(2x2)','99.03%'])
print(x)

+-----+-----+-----+-----+-----+
| No.Convolution_Layers | Kernel_Sizes | No.of_Max_Pooling_Layers | Pool_Size | Test_Accuracy |
+-----+-----+-----+-----+-----+
| 3 | 98.63% | | (3x3) | | 1 | | (2x2) |
| 4 | 99.35% | | (3x3) | | 2 | | (2x2) |
| 7 | 99.03% | | (3x3) | | 1 | | (2x2) |
+-----+-----+-----+-----+-----+

```